

Mixture of Experts (MoEs) (8x7B)

What is Mixture of Experts?

Mixture of Experts (MoE) is a technique that uses many different sub-models (or “experts”) to improve the quality of LLMs.

Two main components define a MoE:

- **Experts** - Each FFNN layer now has a set of “experts” of which a subset can be chosen. These “experts” are typically FFNNs themselves.
- **Router or gate network** - Determines which tokens are sent to which experts.
- MoE is supposed to replace; the *dense layers*.
- Each layer is composed of 8 feed forward blocks known as experts

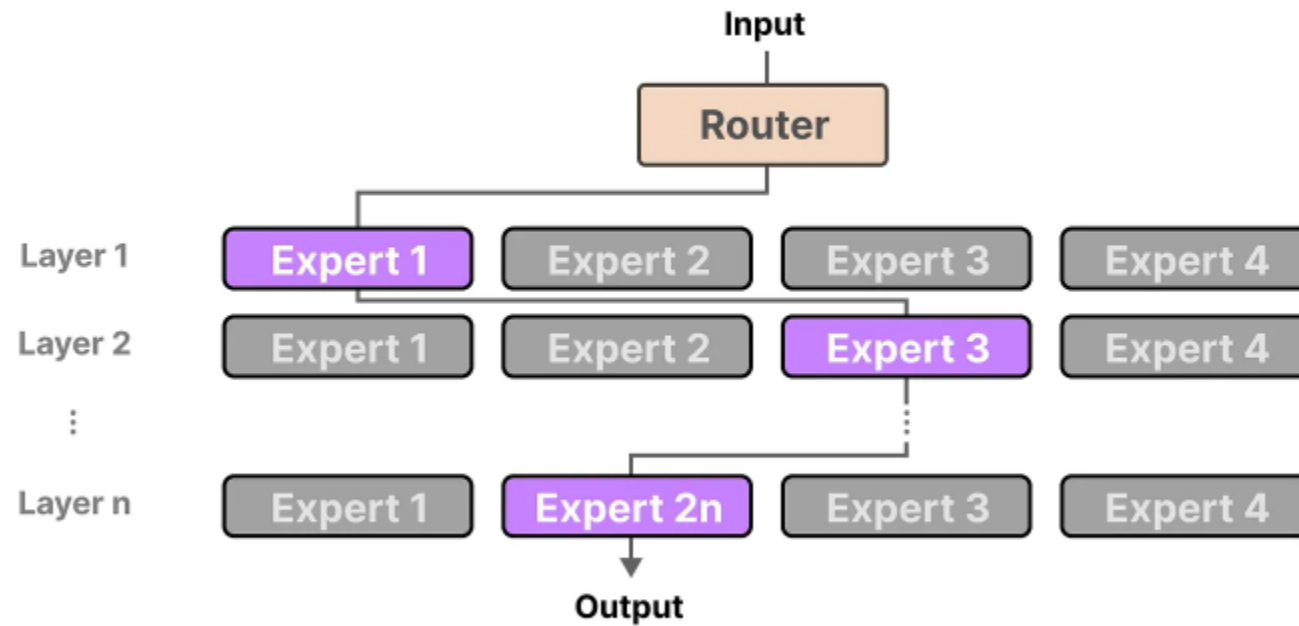
In each layer of an LLM with an MoE, we find (somewhat specialized) experts:



Know that an “expert” is not specialized in a specific domain like “Psychology” or “Biology”. At most, it learns syntactic information on a word level instead:

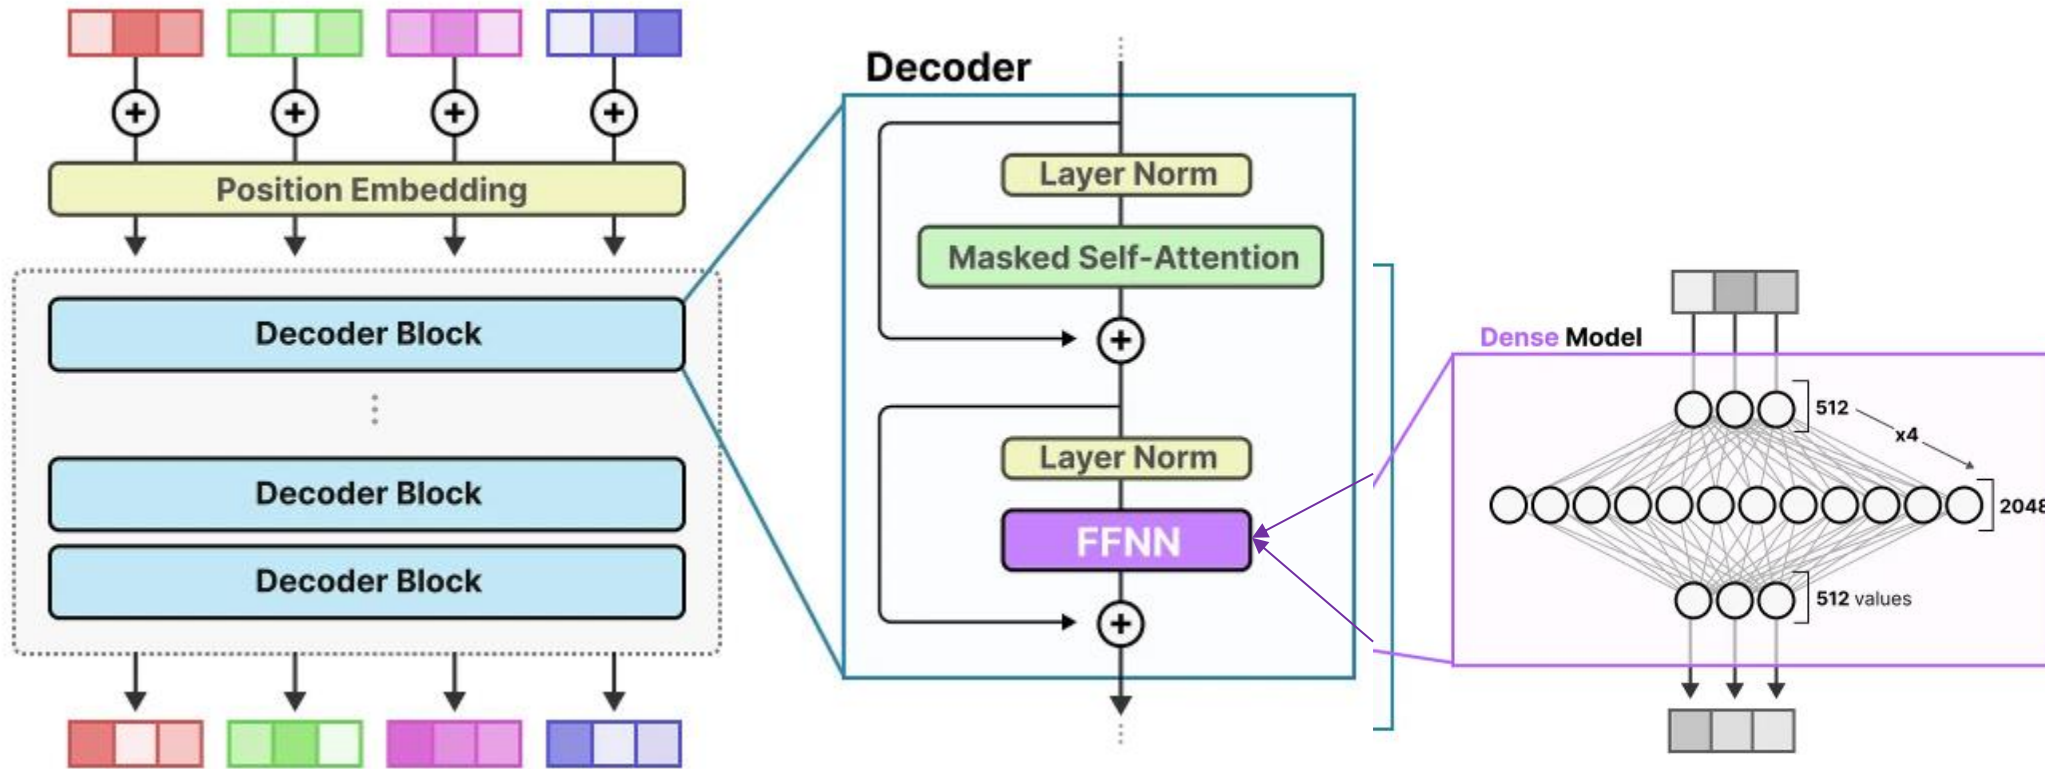


The *router* (gate network) selects the expert(s) best suited for a given input:



Each expert is not an entire LLM but a submodel part of an LLM's architecture.

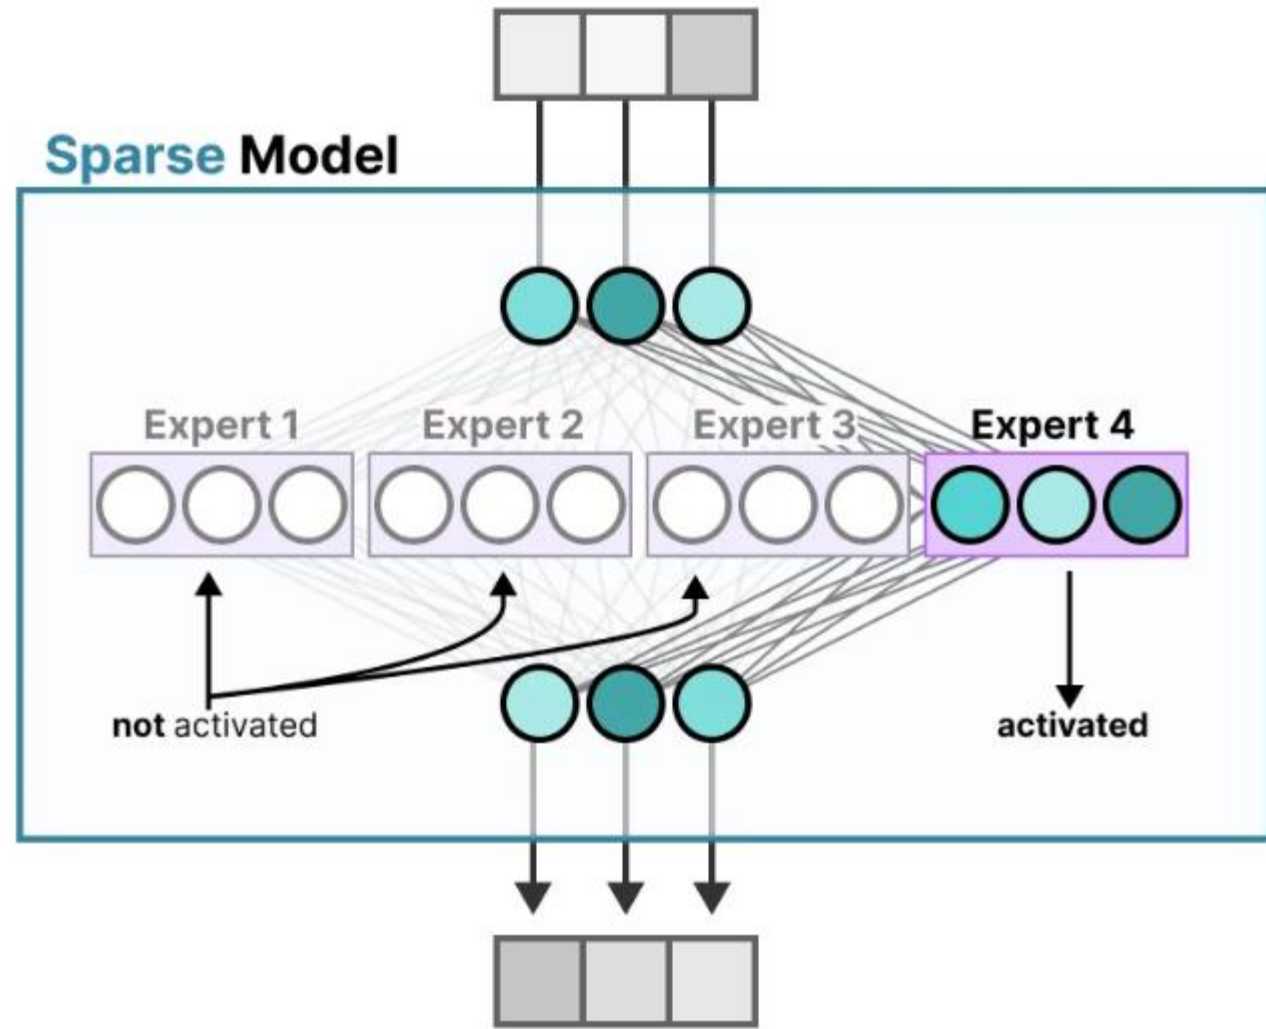
Recall Transformers Arch & FFNN



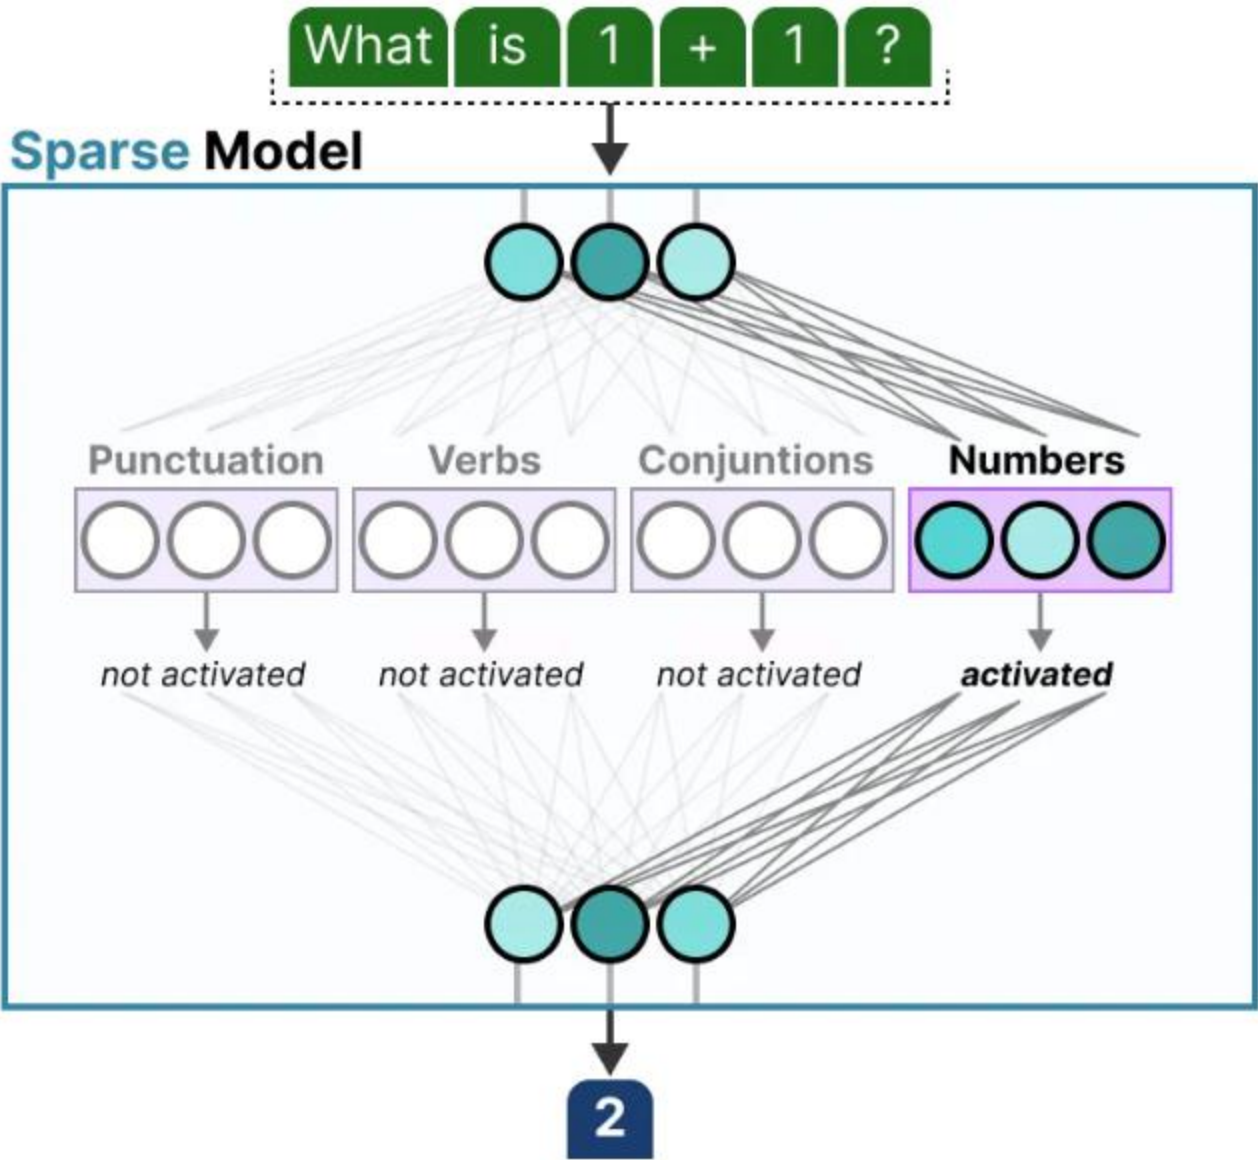
The FFNN in a traditional Transformer is called a **dense** model since all parameters (its weights and biases) are activated.

Sparse Models

- In contrast, **sparse** models only activate a portion of their total parameters and are closely related to Mixture of Experts.
- To illustrate, we can chop up our dense model into pieces (so-called experts), retrain it, and only activate a subset of experts at a given time.
- The underlying idea is that each expert learns different information during training.
- Then, when running inference, only specific experts are used as they are most relevant for a given task.



When asked a question, we can select the expert best suited for a given task.
experts learn more fine-grained information than entire domains



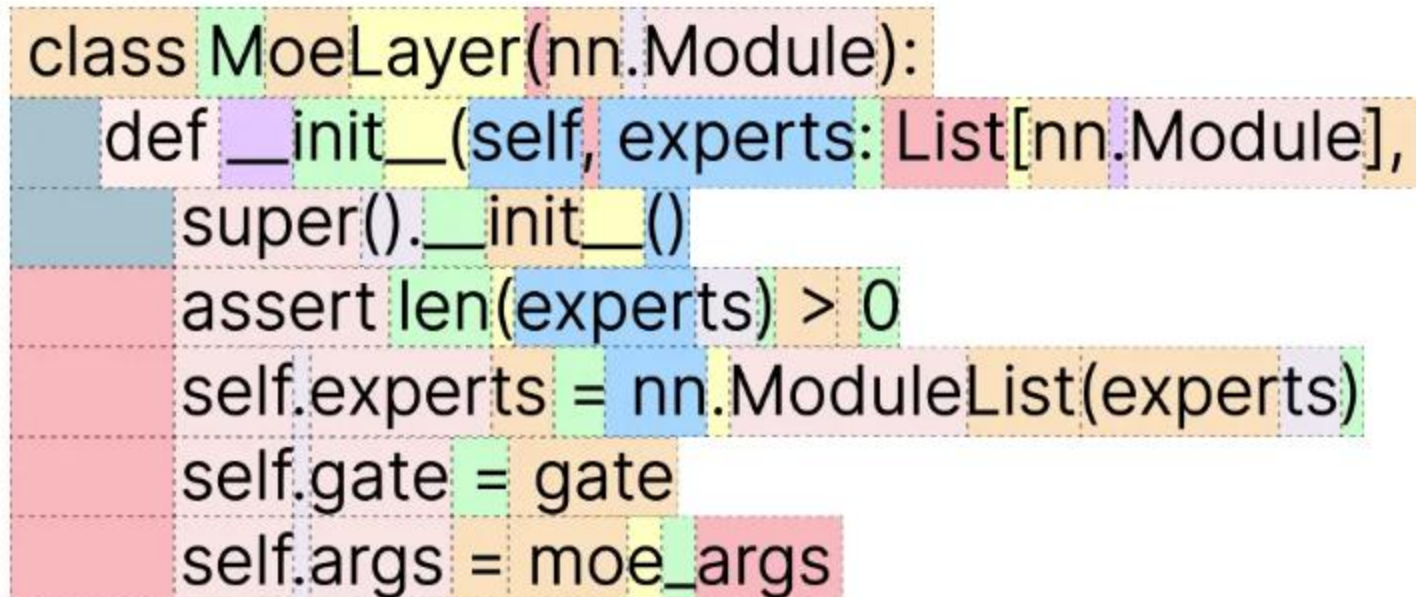
Expert specialization	Expert position	Routed tokens
Punctuation	Layer 2	, , , , , , , - , , , , , , ,) .)
	Layer 6	, , , , , : . : , & , & & ? & - , ? , , , .
Conjunctions and articles	Layer 3	The the the the the the the the The...
	Layer 6	a and and and and and and and or and ...
Verbs	Layer 1	died falling identified fell closed left posted lost felt left said read miss place struggling falling signed died...
Visual descriptions <i>color, spatial position</i>	Layer 0	her over her know dark upper dark outer center upper blue inner yellow raw mama bright bright over open your dark blue
Counting and numbers <i>written and numerical forms</i>	Layer 1	after 37 19. 6. 27 I I Seven 25 4, 54 I two dead we Some 2012 who we few lower

Expert specialization of an encoder model in the ST-MoE paper.

Experts in decoder models, however, do not seem to have the same type of specialization. That does not mean though that all experts are equal.

A great example can be found in the [Mixtral 8x7B paper](#) where each token is colored with the first expert choice.

A great example can be found in the [Mixtral 8x7B paper](#) where each token is colored with the first expert choice.

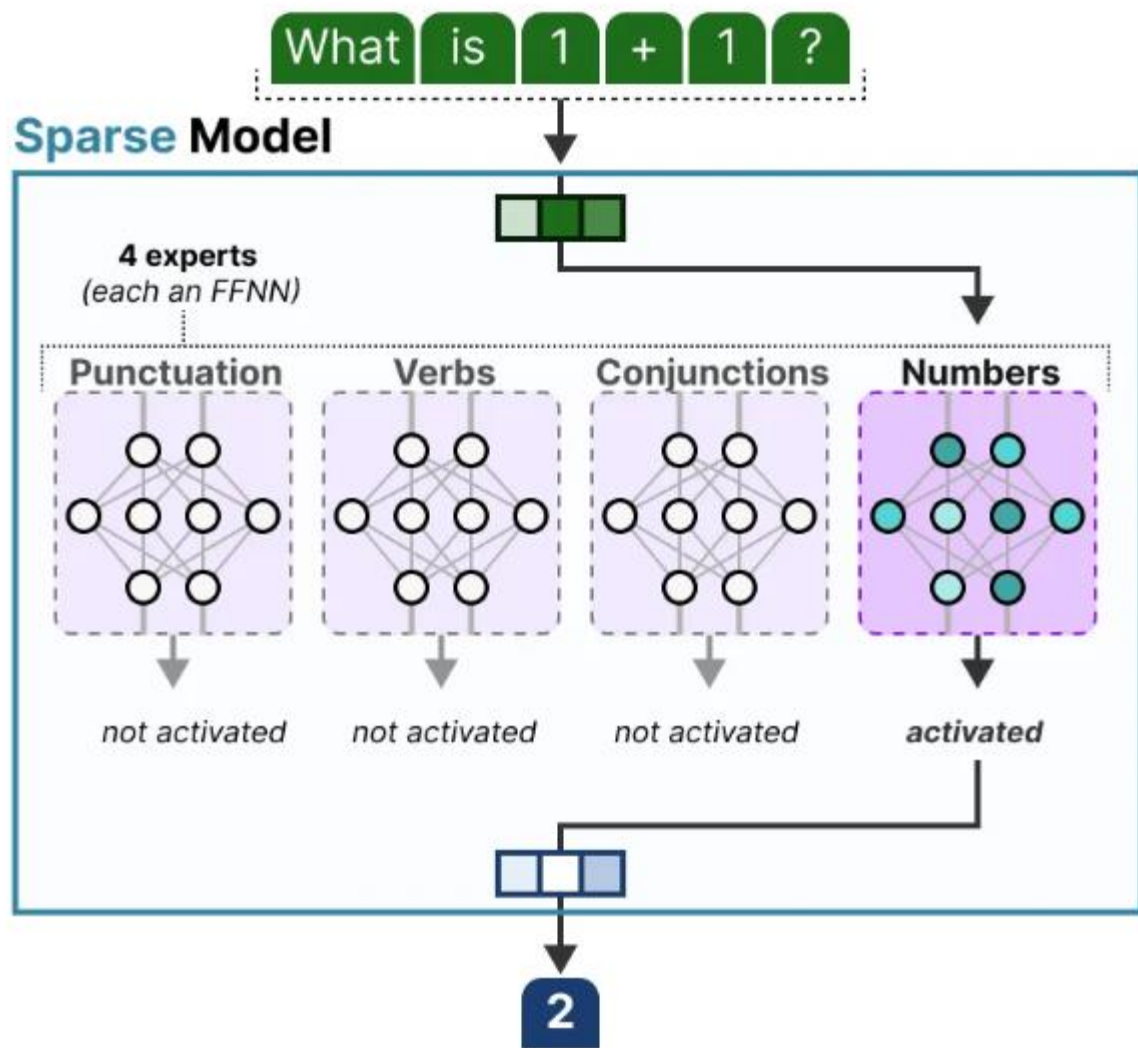


```
class MoeLayer(nn.Module):
    def __init__(self, experts: List[nn.Module],
                 super().__init__())
    assert len(experts) > 0
    self.experts = nn.ModuleList(experts)
    self.gate = gate
    self.args = moe_args
```

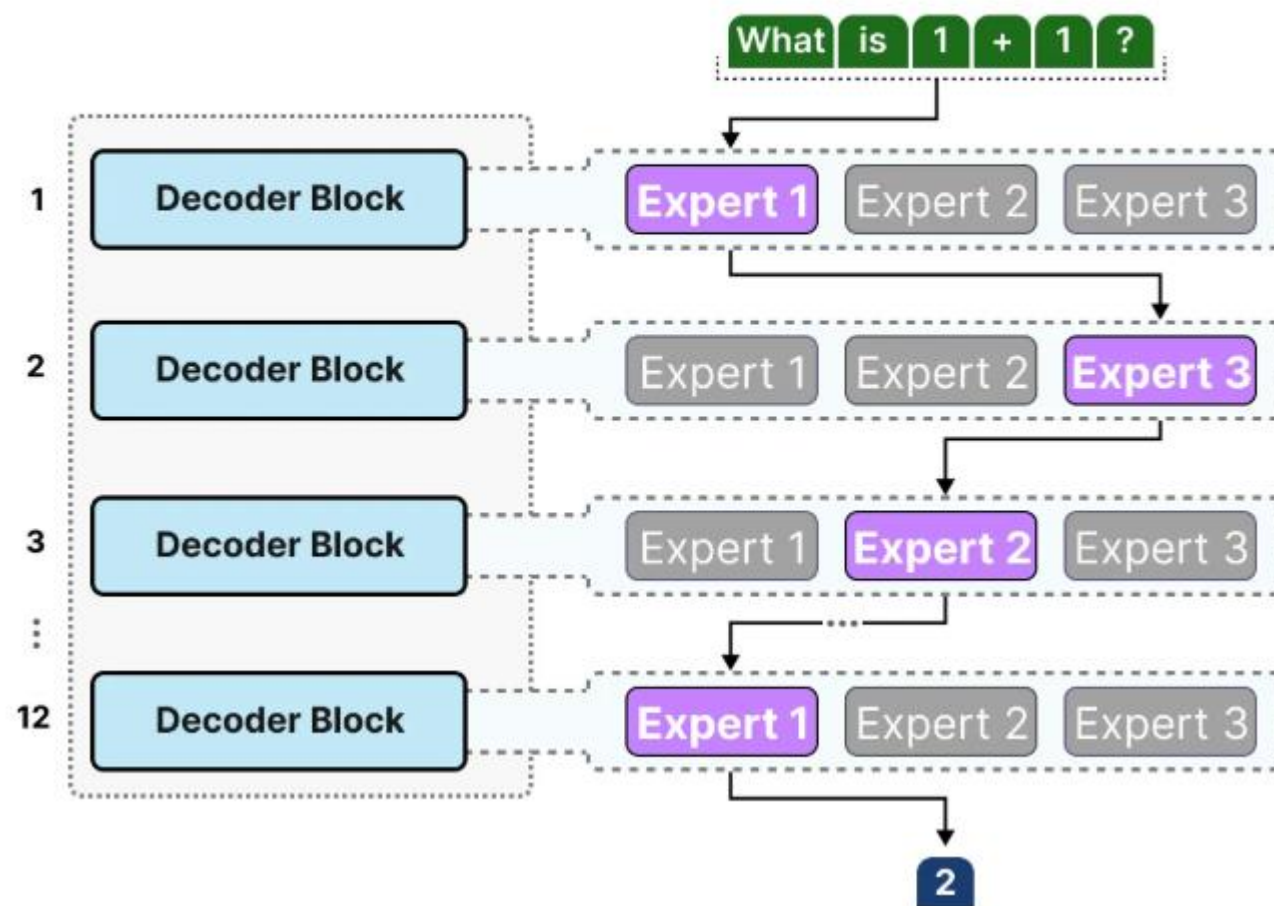
The image shows a Python code snippet for a `MoeLayer` class. Each token in the code is enclosed in a colored box, representing the expert chosen for that token. The colors are: orange for `class`, `MoeLayer`, `nn.Module`, `nn.ModuleList`, `experts`, `gate`, and `moe_args`; green for `__init__` (twice); blue for `self`, `experts` (twice), and `nn.ModuleList`; red for `assert`, `>`, and `args`; light blue for `def`; and grey for `super()`. The boxes are arranged in a grid-like fashion, with some tokens spanning multiple lines.

- Experts don't focus on specific domains (e.g., NLP, math, code) but rather on token types (e.g., function names, variables, operators).
- This visual also demonstrates that experts tend to focus on syntax rather than a specific domain.
- Thus, although decoder experts do not seem to have a specialism they do seem to be used consistently for certain types of tokens.


MoE Architecture



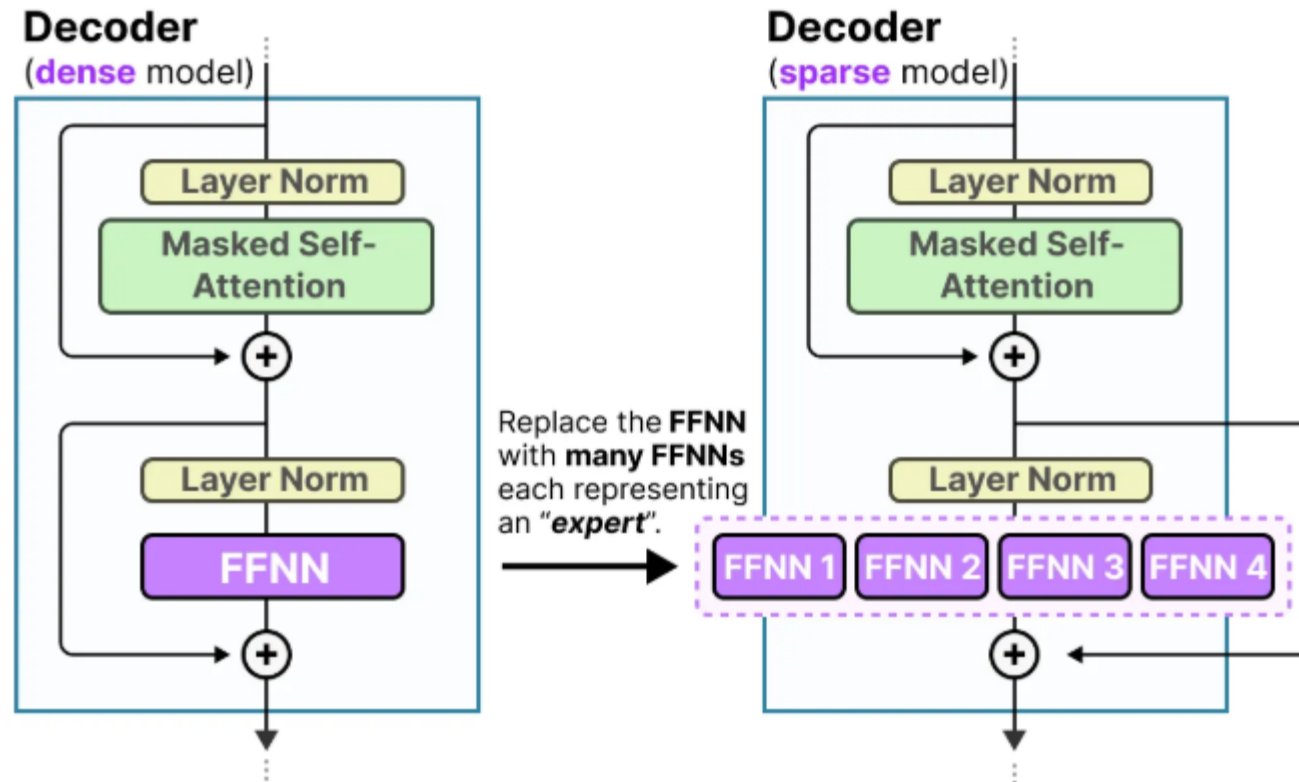
Since most LLMs have several decoder blocks, a given text will pass through multiple experts before the text is generated:



The chosen experts likely differ between tokens which results in different “paths” being taken:



If we update our visualization of the decoder block, it would now contain more FFNNs (one for each expert) instead:



The decoder block now has multiple FFNNs (each an “expert”) that it can use during inference.

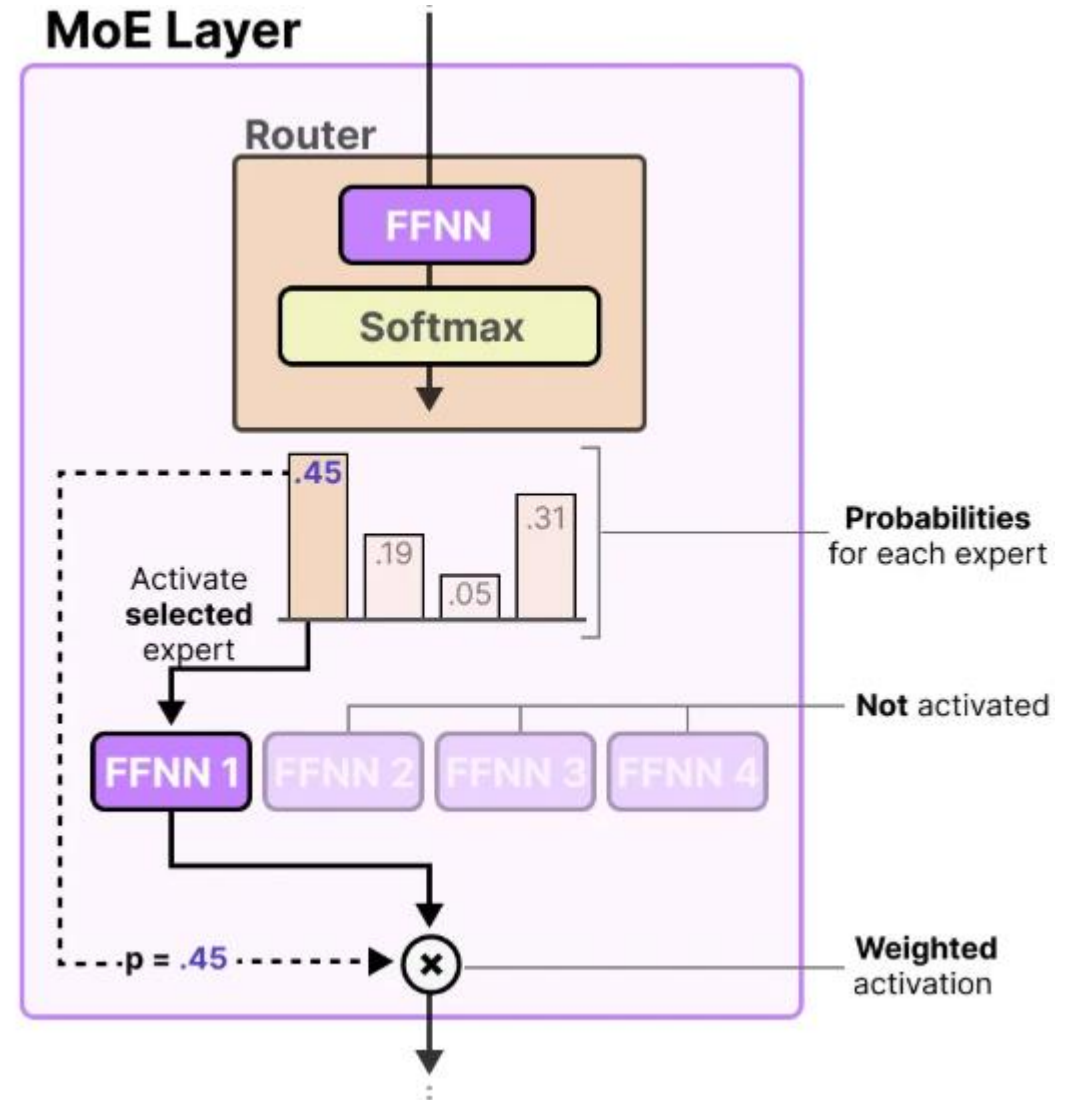
The Routing Mechanism

Now that we have a set of experts, how does the model know which experts to use?

Just before the experts, a **router** (also called a **gate network**) is added which is trained to choose which expert to choose for a given token.

- The **router** (or **gate network**) is also an FFNN and is used to **choose the expert** based on a particular **input**.
- It outputs probabilities which it uses to select the best matching expert.

The expert layer returns the **output** of the selected expert **multiplied** by the **gate value** (selection probabilities).

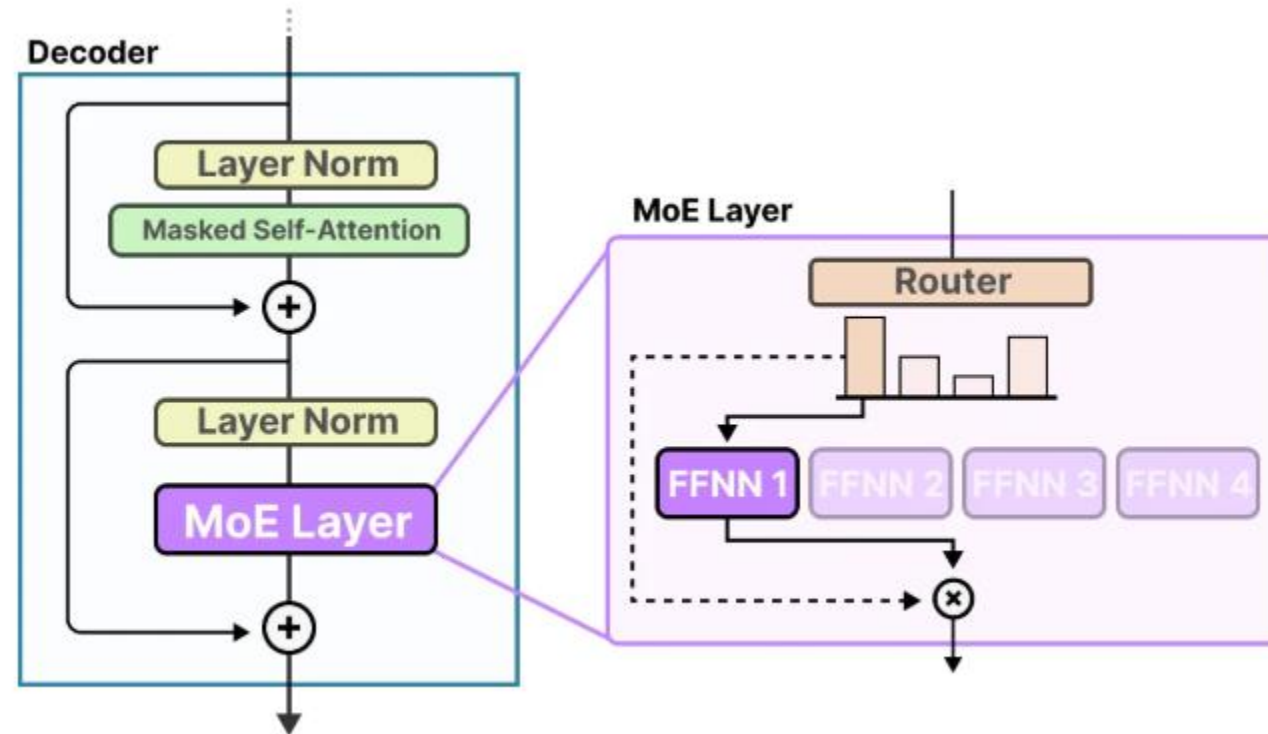


Mistral 8x7B : the gating/routing function

- The gating function is just a linear layer (**in_features=4096, out_feature=8, bias=False**) that's trained along with the rest of the model.
For each token embedding, it produces 8 logits, which indicate which expert to select.

```
Mistral 8x7B → self.feed_forward: nn.Module
                if args.moe is not None:
                    self.feed_forward = MoeLayer(
                        experts=[FeedForward(args=args) for _ in range(args.moe.num_experts)],
                        gate=nn.Linear(args.dim, args.moe.num_experts, bias=False),
                        moe_args=args.moe,
                    )
                else:
Mistral 7B → self.feed_forward = FeedForward(args=args)
```

- The expert layer returns the output of the selected expert multiplied by the gate value (selection probabilities).
- The router together with the experts (of which only a few are selected) makes up the **MoE Layer**:




- A given MoE layer comes in two sizes, either a *sparse* or a **dense** mixture of experts.
- Both use a router to select experts but a **Sparse MoE only selects a few** whereas a **Dense MoE selects them all** but potentially in different distributions.

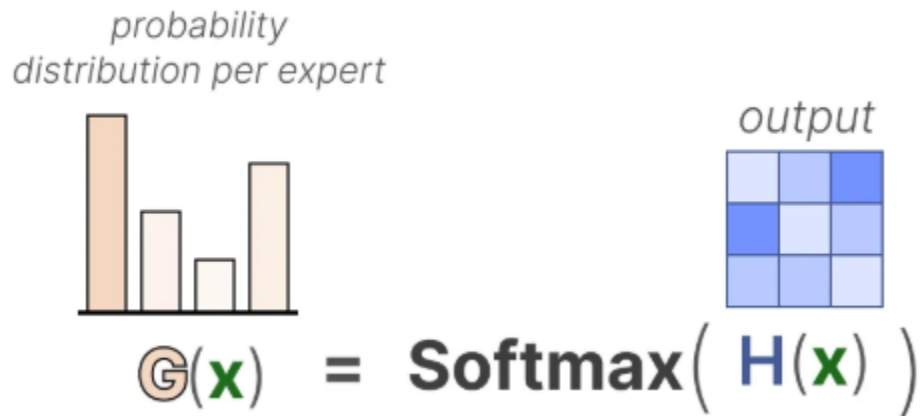
Selection of Experts

The gating network is arguably the most important component of any MoE as it not only decides which experts to choose during *inference* but also *training*.


In its most basic form, we multiply the input (\mathbf{x}) by the router weight matrix (\mathbf{W}):



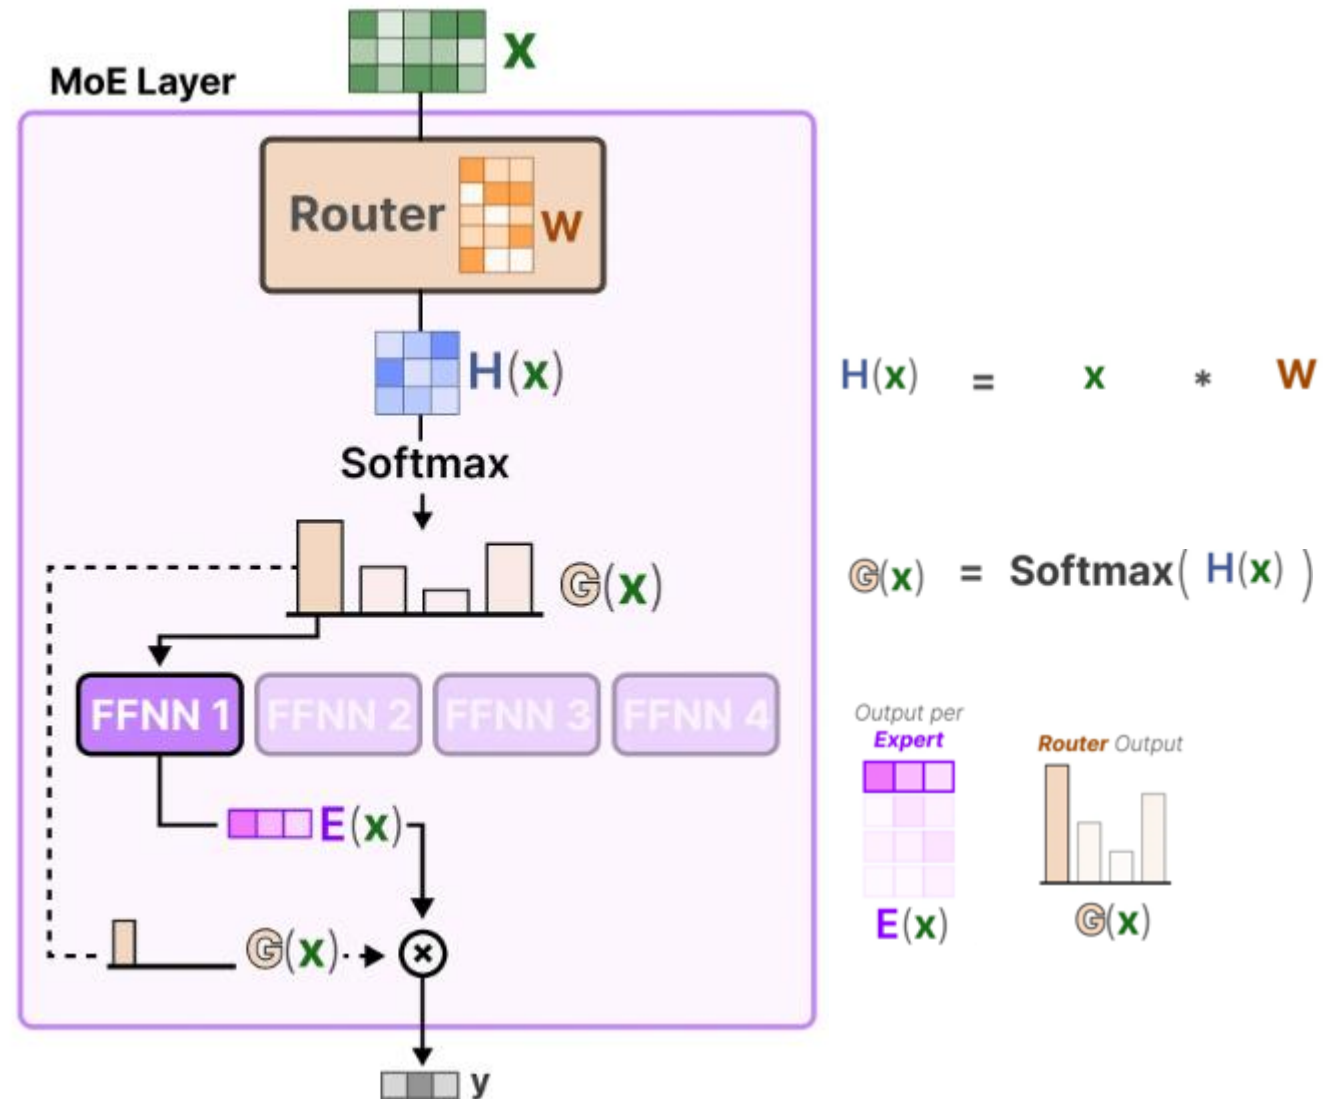
Then, we apply a **SoftMax** on the output to create a probability distribution $G(\mathbf{x})$ per expert:



The router uses this probability distribution to choose the **best matching expert** for a given input.



Let's put everything together and explore how the input flows through the router and experts:

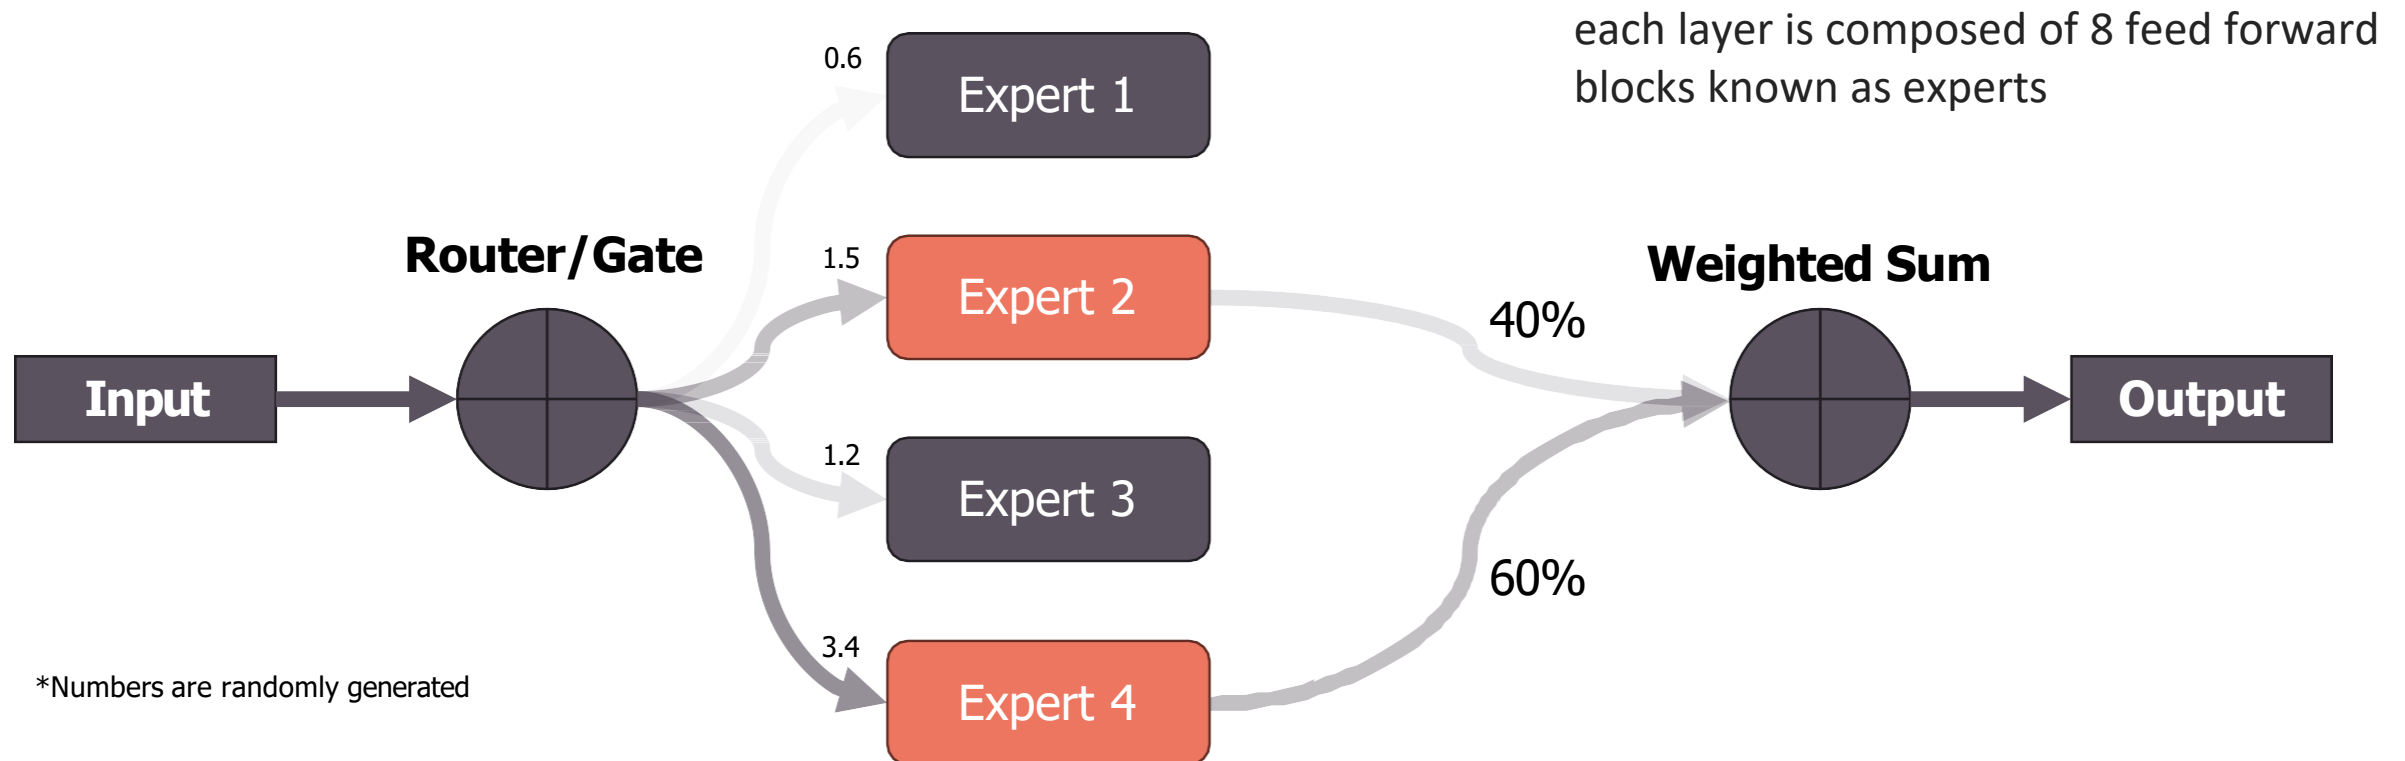


Mixture of Experts: an introduction

Mixture of Experts is an ensemble technique, in which we have multiple “expert” models, each trained on a subset of the data, such that each model specializes on it and then the output of the experts are combined (usually a weighted sum or by averaging) to produce one single output.

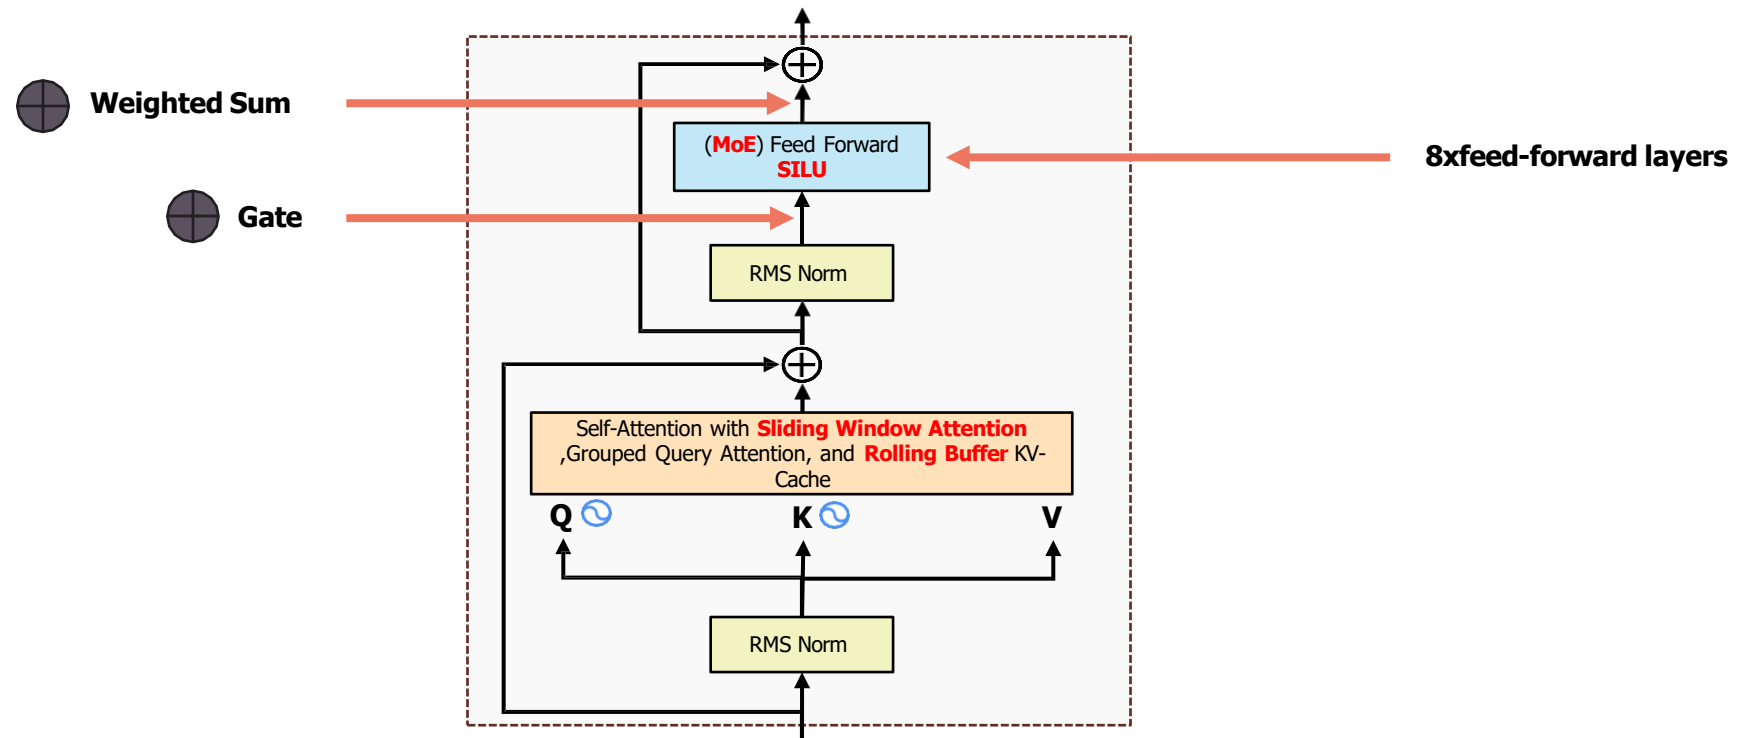
In the case of **Mistral 8x7B**, we talk about **Sparse Mixture of Experts (SMoE)**, because only 2 out of 8 experts are used for every token.

The gate produces **logits** which are used to select the top-k experts. The top-k logits are then run through a **softmax** to produce weights.



Mistral 8x7B: expert feed-forward layers

- In the case of Mistral 8x7B, the experts are the Feed-Forward layers present at every Encoder layer. Each Encoder layer is comprised of a single Self-Attention mechanism, followed by a mixture of experts of 8 FFN. The gate function selects the top 2 experts for each incoming token. The output is combined with a weighted sum.
- This allows to increase the parameters of the model, but without impacting the computation time, since the input will only pass through the top 2 experts, so the intermediate matrix multiplications will be performed only on the selected experts.

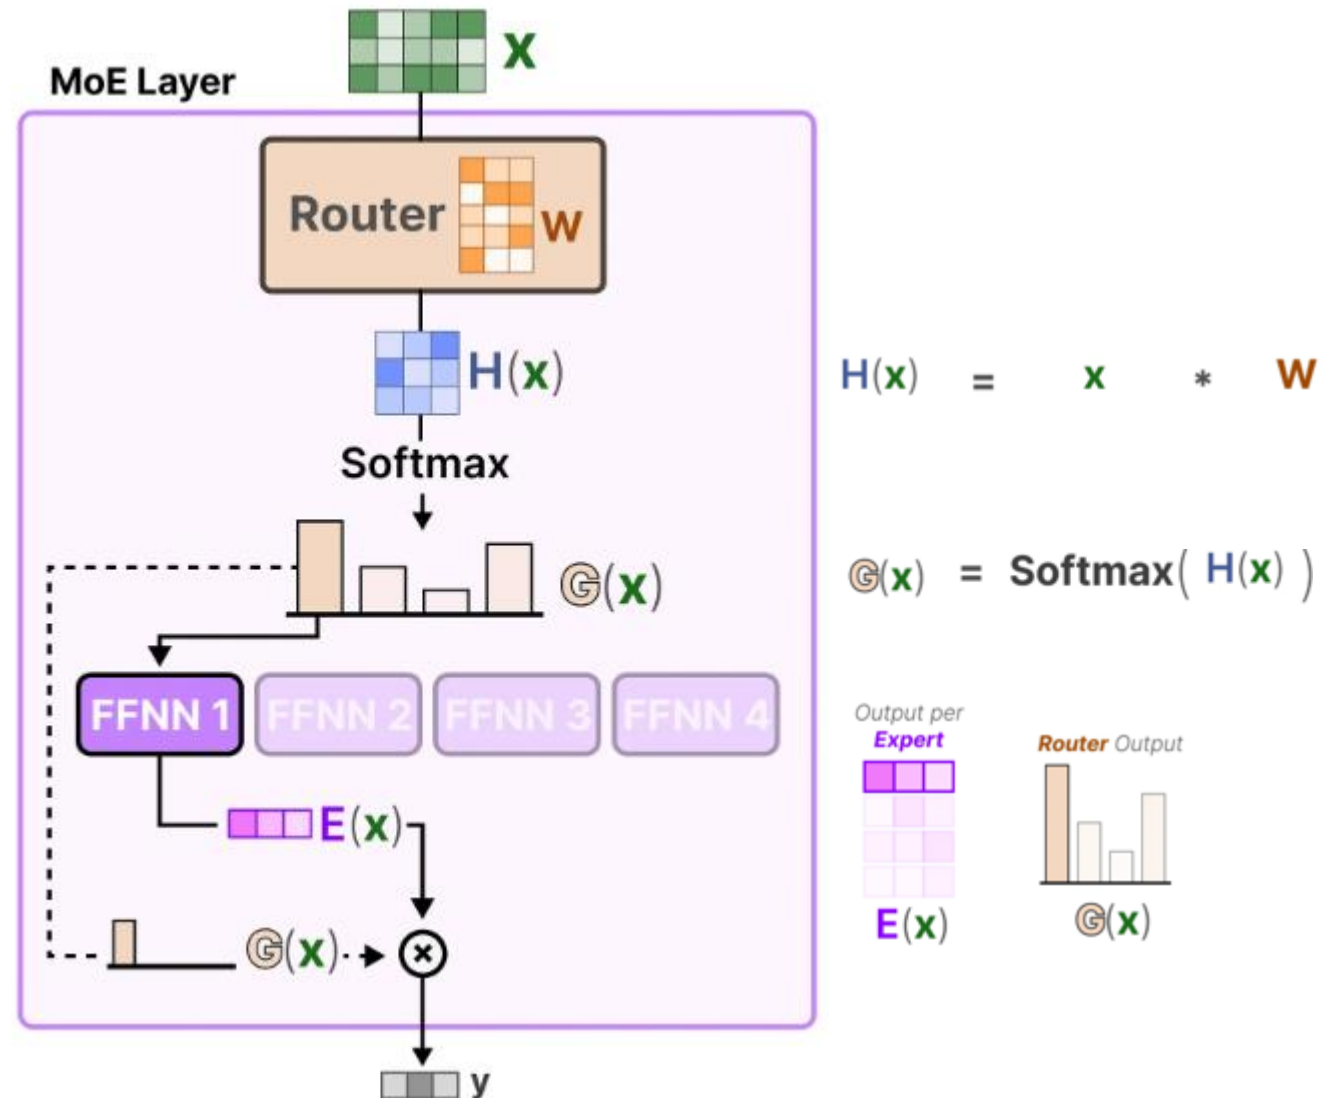


Mixture of Experts: why we apply the softmax after selecting the top-k experts?

- If we apply the softmax directly to the output of the gate function, **this will result in a “probability distribution” (weights that sum up to 1) over all the experts.**
- But since we are only going to use the top-k of them, we want a “probability distribution” over only the selected experts.
- This also makes it easier to compare two models trained on different number of experts, since the sum of the weights applied to the output will always be 1 independently on the number of experts chosen by the gate function.
- Applying softmax after selection has some issues:
- If an expert receives a very small probability due to normalization over all experts, it might still get selected, even though it isn't the best candidate.
- Also, models trained with **different numbers of experts** would behave differently because the softmax probabilities would be normalized over different ranges.

Let's put everything together and explore how the input flows through the router and experts:

Complexity of Routing:
What if the same experts are chosen over others?

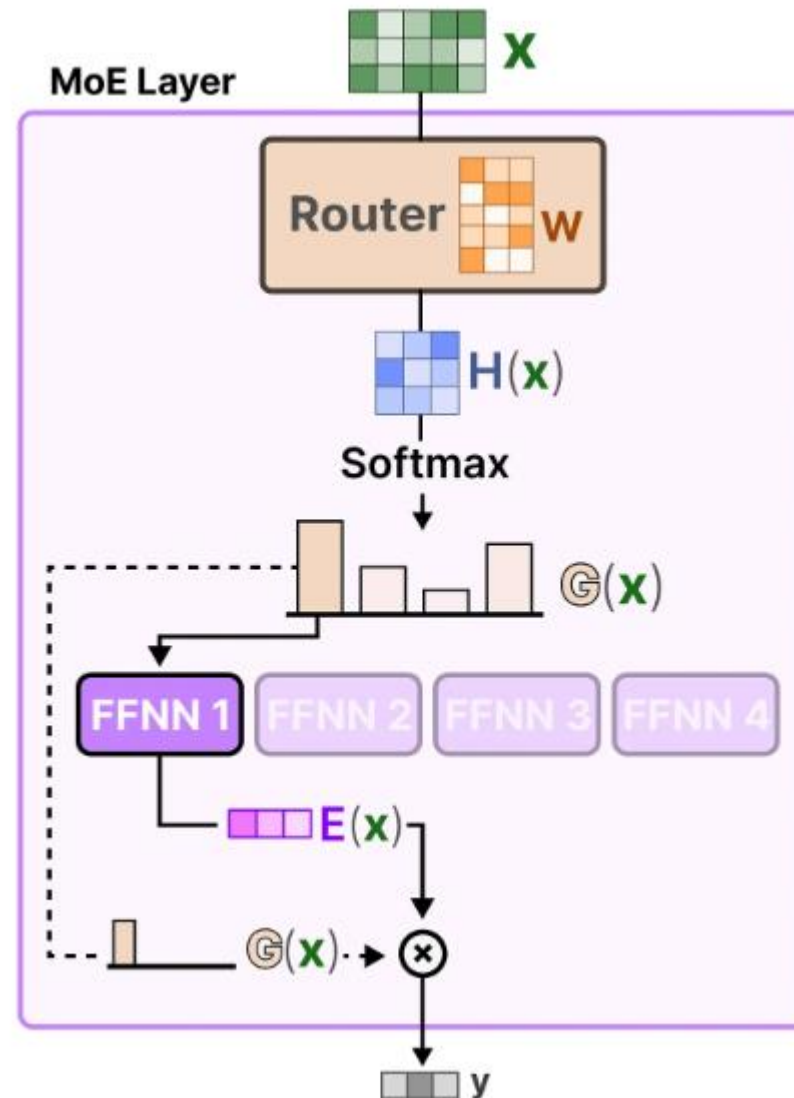


Let's put everything together and explore how the input flows through the router and experts:

Complexity of Routing:

What if the same experts are chosen over others?


- Not only will there be an **uneven distribution of experts** chosen, but some experts will hardly be trained at all.
- This results in issues during both training and inference.
- Instead, we want **equal importance** among experts during training and inference, which we call **load balancing**. In a way, it's to prevent overfitting on the same experts.



$$H(X) = X * W$$

$$G(X) = \text{Softmax}(H(X))$$

Output per Expert



Router Output

